

# CSCI596 Assignment 1—Complexity, Flop/s and Message Passing Interface

Due: September 12 (Wed), 2018, 11:59 pm

## Part I. Computational Complexity and Flop/s Performance


In this part, you will perform small hands-on exercises to get a feel for computational complexity and flop/s (floating-point operations per second) performance of a computer. Submit your answers to the following two questions.

### I-1. Measuring Computational Complexity

Download the data file, `MDtime.out`, from the course home page for assignment 1. In the two-column file, the left column is the number of atoms,  $N$ , simulated by the `md.c` program, whereas the right column is the corresponding running time,  $T$ , of the program in seconds. Make a log-log plot of  $T$  vs.  $N$ . Perform linear fit of  $\log T$  vs.  $\log N$ , i.e.,  $\log T = \alpha \log N + \beta$ , where  $\alpha$  and  $\beta$  are fitting parameters. Note that the coefficient  $\alpha$  signifies the power with which the runtime scales as a function of problem size  $N$ :  $T \propto N^\alpha$ . **Submit the plot and your fit of  $\alpha$ .**

### I-2. Theoretical Flop/s Performance

Suppose you have a quadcore processor (a processor equipped with 4 processing cores) operating at a clock speed of 3.0 GHz (i.e., clock ticks  $3 \times 10^9$  times per second), in which each core can operate 1 multiplication and 1 addition operations per clock cycle (e.g., using a special fused multiply-add (FMA) circuit). Assume that each multiply or add operation is performed on vector registers, each holding 4 double-precision operands (Fig. 1).<sup>\*</sup> **What is the theoretical peak performance of your computer in Gflop/s (gigaflop/s or  $10^9$  flops)?**



**Fig. 1:** Vector multiplier (VMUL) loads data from two vector registers, R1 and R3, each holding 4 double-precision numbers, concurrently performs 4 multiplications, and stores the results on vector register R3.

## Part II. Implementing Your Own Global Summation with Message Passing Interface

In this part, you will write your own global summation program (equivalent to `MPI_Allreduce`) using `MPI_Send` and `MPI_Recv`. Your program should run on  $P = 2^l$  processors ( $l = 0, 1, \dots$ ). Each process contributes a partial value, and at the end, all the processes will have the globally summed value of these partial contributions.

Your program will use a communication structure called butterfly, which is structured as a series of pairwise exchanges (see the figure below where messages are denoted by arrows). This structure allows a global reduction among  $P$  processes to be performed in  $\log_2 P$  steps.

<sup>\*</sup> See FLOPS in Wikipedia [<https://en.wikipedia.org/wiki/FLOPS>].

$$\begin{aligned}
& a000 + a001 + a010 + a011 + a100 + a101 + a110 + a111 \\
& = ((a000 + a001) + (a010 + a011)) + ((a100 + a101) + (a110 + a111))
\end{aligned}$$

At each level  $l$ , a process exchanges messages with a partner whose rank differs only at the  $l$ -th bit position in the binary representation (Fig. 2).




Fig. 2: Butterfly network used in hypercube algorithms.

### HYPERCUBE TEMPLATE

We can use the following template to perform a global reduction using any associative operator  $OP$  (such as multiplication or maximum),  $(a \text{ } OP \text{ } b) \text{ } OP \text{ } c = a \text{ } OP \text{ } (b \text{ } OP \text{ } c)$ .

---

```

procedure hypercube(myid, input, logP, output)
begin
  mydone := input;
  for l := 0 to logP-1 do
    begin
      partner := myid XOR 2l;
      send mydone to partner;
      receive hisdone from partner;
      mydone = mydone OP hisdone
    end
  output := mydone
end

```

---

### USE OF BITWISE LOGICAL XOR

Note that

$$\begin{aligned}
0 \text{ XOR } 0 &= 0 \text{ XOR } 1 = 1; \\
0 \text{ XOR } 1 &= 1 \text{ XOR } 0 = 1.
\end{aligned}$$

so that  $a \text{ XOR } 1$  flips the bit  $a$ , *i.e.*,

$$\begin{aligned}
a \text{ XOR } 1 &= \bar{a} \\
a \text{ XOR } 0 &= a
\end{aligned}$$

where  $\bar{a}$  is the complement of  $a$  ( $\bar{a} = 1 - a$  for  $a = 0$  or  $1$ ). In particular,  $\text{myid} \text{ XOR } 2^l$  reverses the  $l$ -th bit of the rank of this process,  $\text{myid}$ :

$$\text{abcdefg} \text{ XOR } 0000100 = \text{abcd} \bar{e} \text{fg}$$

Note that the XOR operator is  $\wedge$  (caret symbol) in the C programming language.

### ASSIGNMENT

Complete the following program by implementing the function, `global_sum`, using `MPI_Send` and `MPI_Recv` functions and the hypercube template given above.

*Submit the source code as well as the printout from a test run on 4 processors and that on 8 processors.*

---

```
#include "mpi.h"
#include <stdio.h>

int nprocs; /* Number of processors */
int myid;   /* My rank */

double global_sum(double partial) {
    /* Implement your own global summation here */
}

int main(int argc, char *argv[]) {
    double partial, sum, avg;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);

    partial = (double) myid;
    printf("Node %d has %le\n", myid, partial);

    sum = global_sum(partial);

    if (myid == 0) {
        avg = sum/nprocs;
        printf("Global average = %le\n", avg);
    }

    MPI_Finalize();

    return 0;
}
```

---