

Assignment06

From Siqi Liang (liangsiq@usc.edu)

Part I

1. Submit the code

```
/*-----
Program pdf0.c computes a pair distribution function for n atoms
given the 3D coordinates of the atoms.
-----*/

#include <stdio.h>
#include <math.h>
#include <time.h>
#include <stdlib.h>

#define NHBIN 2000 // Histogram size

float al[3]; // Simulation box lengths
int n; // Number of atoms
float *r; // Atomic position array
FILE *fp;

__device__ float d_SignR(float v, float x) {if (x > 0) return v; else return -v;} // jd

__constant__ float DALTH[3];
__constant__ int DN;
__constant__ float DDRH;

float SignR(float v, float x) {if (x > 0) return v; else return -v;}

//jd
__global__ void gpu_histogram_kernel(float *r, float *nhis)
{
    int i, j, a, ih;
    float rij, dr;

    int iBlockBegin = (DN/gridDim.x)*blockIdx.x;
    int iBlockEnd = min((DN/gridDim.x)*(blockIdx.x+1), DN);
    int jBlockBegin = (DN/gridDim.y)*blockIdx.y;
```

```

int jBlockEnd = min((DN/gridDim.y)*(blockIdx.y+1),DN);
for (i=iBlockBegin+threadIdx.x; i<iBlockEnd; i+=blockDim.x)
{
for (j=jBlockBegin+threadIdx.y; j<jBlockEnd; j+=blockDim.y)
{
    if (i<j) { // Process (i,j) atom pair
        rij = 0.0;
        for (a=0; a<3; a++) {
            dr = r[3*i+a]-r[3*j+a];
            /* Periodic boundary condition */
            dr = dr - d_SignR(DALTH[a],dr-DALTH[a]) -
d_SignR(DALTH[a],dr+DALTH[a]);
            rij += dr*dr;
        }
        rij = sqrt(rij); /* Pair distance */
        ih = rij/DDRH;
        atomicAdd(&nhis[ih],1.0);
    } // end if i<j
} // end for j
} // end for i
}
//

/*-----*/
void histogram() {
/*-----
Constructs a histogram NHIS for atomic-pair distribution.
-----*/

float alth[3];
float* nhis; // Histogram array
float rhmax,drh,density,gr;
int a,ih;

float* dev_r; // Atomic positions
float* dev_nhis; // Histogram

/* Half the simulation box size */
for (a=0; a<3; a++) alth[a] = 0.5*al[a];
/* Max. pair distance RHMAX & histogram bin size DRH */
rhmax = sqrt(alth[0]*alth[0]+alth[1]*alth[1]+alth[2]*alth[2]);
drh = rhmax/NHBIN; // Histogram bin size

nhis = (float*)malloc(sizeof(float)*NHBIN);
//for (ih=0; ih<NHBIN; ih++) nhis[ih] = 0.0; // Reset the histogram

cudaMalloc((void*)&dev_r,sizeof(float)*3*n);

```

```

cudaMalloc((void*)&dev_nhis, sizeof(float)*NHBIN);

cudaMemcpy(dev_r, r, 3*n*sizeof(float), cudaMemcpyHostToDevice);
cudaMemset(dev_nhis, 0.0, NHBIN*sizeof(float));
cudaMemcpyToSymbol(DALTH, alth, sizeof(float)*3, 0, cudaMemcpyHostToDevice);
cudaMemcpyToSymbol(DN, &n, sizeof(int), 0, cudaMemcpyHostToDevice);
cudaMemcpyToSymbol(DDRH, &drh, sizeof(float), 0, cudaMemcpyHostToDevice);

dim3 numBlocks(8, 8, 1);
dim3 threads_per_block(16, 16, 1);
gpu_histogram_kernel<<<numBlocks, threads_per_block>>>(dev_r, dev_nhis); //
Compute dev_nhis on GPU: dev_r[] @ dev_nhis[]

cudaMemcpy(nhis, dev_nhis, NHBIN*sizeof(float), cudaMemcpyDeviceToHost);
cudaFree(dev_r);
cudaFree(dev_nhis);

density = n/(al[0]*al[1]*al[2]);
/* Print out the histogram */
fp = fopen("pdf.d", "w");
for (ih=0; ih<NHBIN; ih++) {
    gr = nhis[ih]/(2*M_PI*pow((ih+0.5)*drh, 2)*drh*density*n);
    fprintf(fp, "%e %e\n", (ih+0.5)*drh, gr);
}
fclose(fp);
free(nhis);
}


int main() {
    int i;
    float cpu1, cpu2;

    /* Read the atomic position data */
    fp = fopen("pos.d", "r");
    fscanf(fp, "%f %f %f", &(al[0]), &(al[1]), &(al[2]));
    fscanf(fp, "%d", &n);
    r = (float*)malloc(sizeof(float)*3*n);
    for (i=0; i<n; i++)
        fscanf(fp, "%f %f %f", &(r[3*i]), &(r[3*i+1]), &(r[3*i+2]));
    fclose(fp);
    /* Compute the histogram */
    cpu1 = ((float) clock())/CLOCKS_PER_SEC;
    histogram();
    cpu2 = ((float) clock())/CLOCKS_PER_SEC;
    printf("Execution time (s) = %le\n", cpu2-cpu1);
}

```

```
    free(r);  
    return 0;  
}
```

2.



Part II

1. Submit code

```
// Hybrid MPI+CUDA computation of Pi  
#include <stdio.h>  
#include <omp.h> //jd  
#include <mpi.h>  
#include <cuda.h>  
  
#define NBIN 10000000 // Number of bins  
#define NUM_DEVICE 2 // jd  
#define NUM_BLOCK 13 // Number of thread blocks
```

```

#define NUM_THREAD 192 // Number of threads per block

// Kernel that executes on the CUDA device
__global__ void cal_pi(float *sum,int nbin,float step,float offset,int
nthreads,int nblocks) {
    int i;
    float x;
    int idx = blockIdx.x*blockDim.x+threadIdx.x; // Sequential thread index
across the blocks
    for (i=idx; i<nbin; i+=nthreads*nblocks) { // Interleaved bin assignment
to threads
        x = offset+(i+0.5)*step;
        sum[idx] += 4.0/(1.0+x*x);
    }
}

int main(int argc,char **argv) {
    int myid,nproc,nbin,tid, mpid;
    float step,offset,pi=0.0,pig;
    dim3 dimGrid(NUM_BLOCK,1,1); // Grid dimensions (only use 1D)
    dim3 dimBlock(NUM_THREAD,1,1); // Block dimensions (only use 1D)
    float *sumHost,*sumDev; // Pointers to host & device arrays
    int dev_used;

    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid); // My MPI rank
    MPI_Comm_size(MPI_COMM_WORLD,&nproc); // Number of MPI processes
    //nbin = NBIN/nproc; // Number of bins per MPI process
    //step = 1.0/(float)(nbin*nproc); // Step size with redefined number of
bins
    //offset = myid*step*nbin; // Quadrature-point offset

    omp_set_num_threads(NUM_DEVICE); // One OpenMP thread per GPU device
    nbin = NBIN/(nproc*NUM_DEVICE); // # of bins per OpenMP thread
    step = 1.0/(float)(nbin*nproc*NUM_DEVICE);
    #pragma omp parallel private( mpid, offset, sumHost, sumDev, tid,
dev_used ) reduction(+:pi)
    {
        mpid = omp_get_thread_num();
        offset = (NUM_DEVICE*myid+mpid)*step*nbin; // Quadrature-point offset
        cudaSetDevice(mpid%2);
        //cudaSetDevice(myid%2);
        size_t size = NUM_BLOCK*NUM_THREAD*sizeof(float); //Array memory size
        sumHost = (float *)malloc(size); // Allocate array on host
        cudaMalloc((void **) &sumDev,size); // Allocate array on device
        cudaMemset(sumDev,0,size); // Reset array in device to 0
    }
}

```

```

        // Calculate on device (call CUDA kernel)
        cal_pi <<<dimGrid,dimBlock>>>
(sumDev,nbin,step,offset,NUM_THREAD,NUM_BLOCK);
        // Retrieve result from device and store it in host array
        cudaMemcpy(sumHost,sumDev,size,cudaMemcpyDeviceToHost);
        // Reduction over CUDA threads
        for(tid=0; tid<NUM_THREAD*NUM_BLOCK; tid++)
            pi += sumHost[tid];
        pi *= step;// race condition solved by reduction(+:pi)

        // CUDA cleanup
        free(sumHost);
        cudaFree(sumDev);
        cudaGetDevice(&dev_used);
        //printf("myid = %d: device used = %d; partial pi =
%f\n",myid,dev_used,pi);
        printf("myid = %d; mpid = %d: device used = %d; partial pi = %f\n",
myid, mpid, dev_used, pi);
    } // END omp parallel jd

    // Reduction over MPI processes
    MPI_Allreduce(&pi,&pig,1,MPI_FLOAT,MPI_SUM,MPI_COMM_WORLD);
    if (myid==0) printf("PI = %f\n",pig);

    MPI_Finalize();
    return 0;
}

```

2.

```

myid = 1; mpid = 1: device used = 1; partial pi = 0.567582
myid = 1; mpid = 0: device used = 0; partial pi = 0.719409
myid = 0; mpid = 1: device used = 1; partial pi = 0.874671
myid = 0; mpid = 0: device used = 0; partial pi = 0.979926
PI = 3.141588

```